


HW3-MATH4323

anthonymcastillo ID:1670011

2022-10-03

Question 1 (a):

```
x1 <- -10:10
x2 <- x1
plot(x1,x2, col = "purple",type = "l")
points(1,1, pch = 19, col = "blue")
points(0,-1,pch = 19,col = "blue")
points(0,1, pch = 19,col = "red")
text(-2,5,"X1 - X2 < 0", col = "red")
text(5,-2,"X1 - X2 > 0", col = "blue")
```




Question 1 (b):

```
x1 <- -10:10
x2 <- -x1
plot(x1,x2, type = "l", col = "magenta")
points(1,1, pch = 19, col = "red")
```

```

points(-1,-1,pch = 19,col = "blue")
points(1,-1, pch = 19,col = "blue")
text(-4,-2,"X_1 + X_2 < 0", col = "blue")
text(4,2,"X_1 + X_2 > 0", col = "red")

```




Question 2 (a):

```


plot(NA, NA, type = "n", xlim = c(-2, -1), ylim = c(-2, 6), asp = 1, xlab =
"X1", ylab = "X2")
symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE)
symbols(c(-1), c(2), circles = c(4), add = TRUE, inches = FALSE)

```




Question 2 (b):

```
plot(NA, NA, type = "n", xlim = c(-2, -1), ylim = c(-2, 6), asp = 1, xlab =  
"X1", ylab = "X2")  
symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE)  
symbols(c(-1), c(2), circles = c(4), add = TRUE, inches = FALSE)  
text(4, 2, ">16", col = "red")  
text(-1, 2, "<4", col = "red")  
text(1, 4, "16<= , >=4", col = "blue")
```



Question 2 (c):

```
plot(NA, NA, type = "n", xlim = c(-2, -1), ylim = c(-2, 6), asp = 1, xlab =
"X1", ylab = "X2")
symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE)
symbols(c(-1), c(2), circles = c(4), add = TRUE, inches = FALSE)
text(4,2,">16", col = "red")
text(-1,2,"<4", col = "red")
text(1,4,"16<= , >=4", col = "blue")
points(0,0,pch=19,col = "blue")
points(-1,1,pch=19,col = "red")
points(2,2,pch=19,col = "blue")
points(3,4,pch=19,col = "red")
```




Question 2 (d):

$(1 + X_1)^2 + (2 - X_2)^2 = 4$ is linear by expanding it $1 + X_1 + X_1^2 + X_2 + X_2^2 = 0$


Question 3 (a):

```
x1 <- c(1,3,4,2,4,4,1,1)
x2 <- c(4,4,3,2,2,4,2,3)
plot.color <- c("blue","red","red","blue","red","red","blue","blue")
plot(x1,x2, col = plot.color, xlim = c(0,5), ylim = c(0,5), pch = 19)
```



Question 3 (b):

```
plot(x1,x2, col = plot.color, xlim = c(0,5), ylim = c(0,5), pch = 19)  
abline(8,-2)
```




equation for hyperplane: $x_1 + x_2 + 8 = 0$

Question 3 (c):

classify to blue if $x_1 + x_2 + 1 < 0$ and classify to red if $x_1 + x_2 + 1 > 0$

Question 3 (d):

```
plot(x1,x2, col = plot.color, xlim = c(0,5), ylim = c(0,5), pch = 19)
abline(8,-2)
abline(6.0,-2.0, lty = 2)
abline(10.0,-2.0, lty = 2)
```



Question 3 (e):


The support vectors are (2,2),(1,4),(4,2),and (3,4)

Question 3 (f):

The 7th obs is (1,2) which is not a support vector which will not change the maximal margin hyperplane

Question 3 (g):


```
plot(x1,x2, col = plot.color, xlim = c(0,5), ylim = c(0,5), pch = 19)
abline(7,-2)
```

equation for hyperplane: $x_1 + x_2 + 7 = 0$

Question 3 (h):

```
plot(x1,x2, col = plot.color, xlim = c(0,5), ylim = c(0,5), pch = 19)
abline(8,-2)
abline(6.0,-2.0, lty = 2)
abline(10.0,-2.0, lty = 2)
points(4,5, pch = 19, col = "blue")
```



Question 4 (a):

```
library(MASS)
library(ISLR)

## Warning: package 'ISLR' was built under R version 4.1.3

library(class)
library(caret)

## Warning: package 'caret' was built under R version 4.1.3
## Loading required package: ggplot2
## Warning: package 'ggplot2' was built under R version 4.1.3
## Loading required package: lattice

newBoston <- data.frame(Boston)
medv.med <- median(Boston$medv)
newBoston$medv01 <- ifelse(Boston$medv > medv.med, yes = 1, no = 0)
newBoston$medv = NULL
```

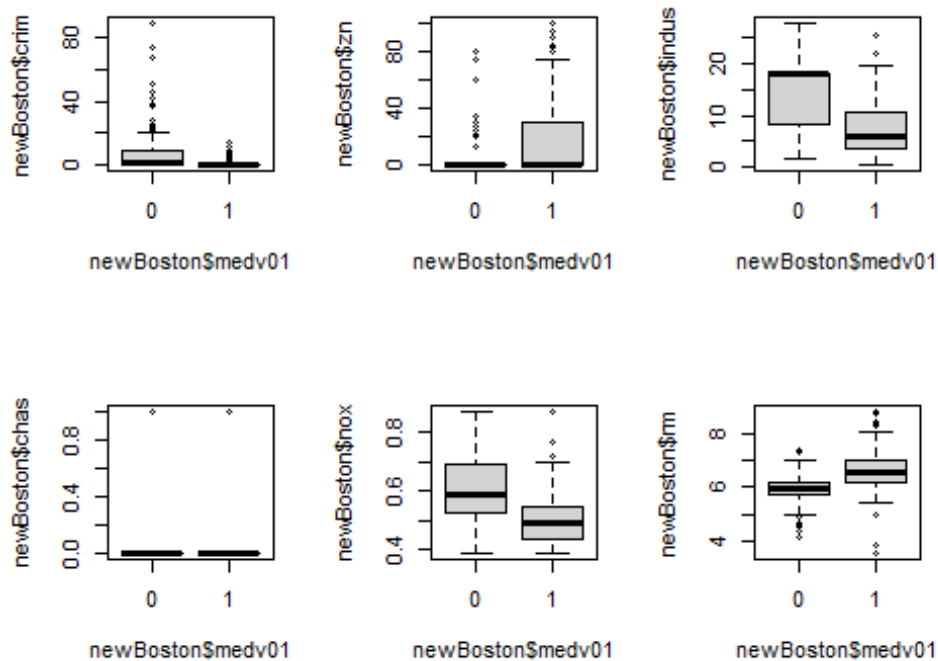
Question 4 (b):

```
cor(newBoston)
```

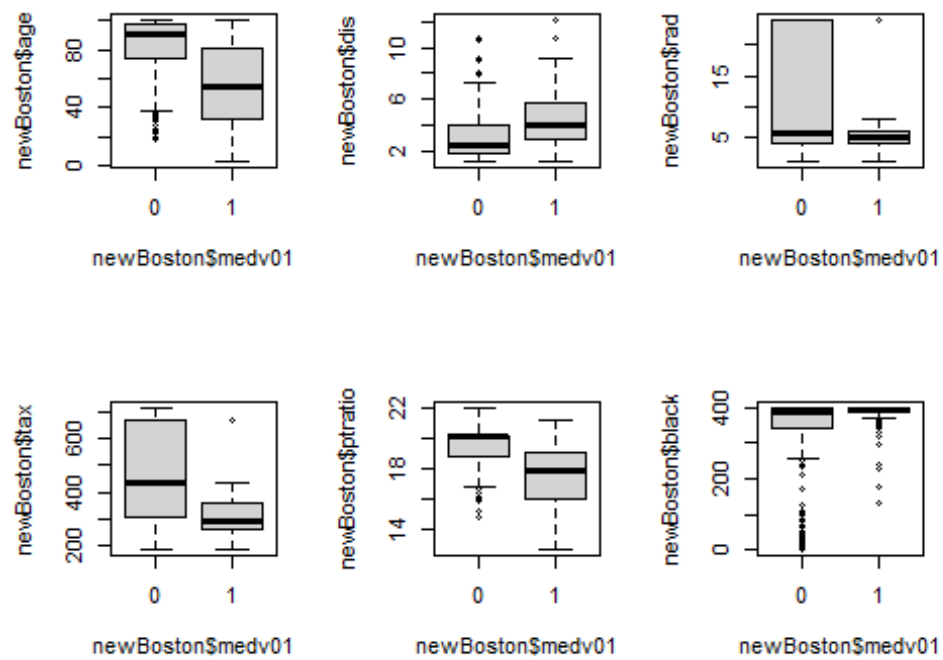
##		crim	zn	indus	chas	nox
## crim	1.00000000	-0.20046922	0.40658341	-0.055891582	0.42097171	
## zn	-0.20046922	1.00000000	-0.53382819	-0.042696719	-0.51660371	
## indus	0.40658341	-0.53382819	1.00000000	0.062938027	0.76365145	
## chas	-0.05589158	-0.04269672	0.06293803	1.000000000	0.09120281	
## nox	0.42097171	-0.51660371	0.76365145	0.091202807	1.00000000	
## rm	-0.21924670	0.31199059	-0.39167585	0.091251225	-0.30218819	
## age	0.35273425	-0.56953734	0.64477851	0.086517774	0.73147010	
## dis	-0.37967009	0.66440822	-0.70802699	-0.099175780	-0.76923011	
## rad	0.62550515	-0.31194783	0.59512927	-0.007368241	0.61144056	
## tax	0.58276431	-0.31456332	0.72076018	-0.035586518	0.66802320	
## ptratio	0.28994558	-0.39167855	0.38324756	-0.121515174	0.18893268	
## black	-0.38506394	0.17552032	-0.35697654	0.048788485	-0.38005064	
## lstat	0.45562148	-0.41299457	0.60379972	-0.053929298	0.59087892	
## medv01	-0.31397528	0.33050237	-0.47164023	0.135646949	-0.45000549	
##	rm	age	dis	rad	tax	
ptratio						
## crim	-0.21924670	0.35273425	-0.37967009	0.625505145	0.58276431	
0.2899456						
## zn	0.31199059	-0.56953734	0.66440822	-0.311947826	-0.31456332	-
0.3916785						
## indus	-0.39167585	0.64477851	-0.70802699	0.595129275	0.72076018	
0.3832476						
## chas	0.09125123	0.08651777	-0.09917578	-0.007368241	-0.03558652	-
0.1215152						
## nox	-0.30218819	0.73147010	-0.76923011	0.611440563	0.66802320	
0.1889327						
## rm	1.00000000	-0.24026493	0.20524621	-0.209846668	-0.29204783	-
0.3555015						
## age	-0.24026493	1.00000000	-0.74788054	0.456022452	0.50645559	
0.2615150						
## dis	0.20524621	-0.74788054	1.00000000	-0.494587930	-0.53443158	-
0.2324705						
## rad	-0.20984667	0.45602245	-0.49458793	1.000000000	0.91022819	
0.4647412						
## tax	-0.29204783	0.50645559	-0.53443158	0.910228189	1.00000000	
0.4608530						
## ptratio	-0.35550149	0.26151501	-0.23247054	0.464741179	0.46085304	
1.0000000						
## black	0.12806864	-0.27353398	0.29151167	-0.444412816	-0.44180801	-
0.1773833						
## lstat	-0.61380827	0.60233853	-0.49699583	0.488676335	0.54399341	
0.3740443						
## medv01	0.50961542	-0.47493222	0.30185475	-0.359152204	-0.45097795	-
0.4601715						
##	black	lstat	medv01			
## crim	-0.38506394	0.4556215	-0.3139753			
## zn	0.17552032	-0.4129946	0.3305024			
## indus	-0.35697654	0.6037997	-0.4716402			
## chas	0.04878848	-0.0539293	0.1356469			

```
## nox      -0.38005064  0.5908789 -0.4500055
## rm       0.12806864 -0.6138083  0.5096154
## age      -0.27353398  0.6023385 -0.4749322
## dis      0.29151167 -0.4969958  0.3018547
## rad      -0.44441282  0.4886763 -0.3591522
## tax      -0.44180801  0.5439934 -0.4509779
## ptratio  -0.17738330  0.3740443 -0.4601715
## black    1.00000000 -0.3660869  0.3029397
## lstat    -0.36608690  1.0000000 -0.6633267
## medv01   0.30293974 -0.6633267  1.0000000
```


```
par(mfrow=c(2,3))
boxplot(newBoston$crim~newBoston$medv01)
boxplot(newBoston$zn~newBoston$medv01)
boxplot(newBoston$indus~newBoston$medv01)
boxplot(newBoston$chas~newBoston$medv01)
boxplot(newBoston$nox~newBoston$medv01)
boxplot(newBoston$rm~newBoston$medv01)
```



```
boxplot(newBoston$age~newBoston$medv01)
boxplot(newBoston$dis~newBoston$medv01)
boxplot(newBoston$rad~newBoston$medv01)
boxplot(newBoston$tax~newBoston$medv01)
boxplot(newBoston$ptratio~newBoston$medv01)
boxplot(newBoston$black~newBoston$medv01)
```



```
boxplot(newBoston$lstat~newBoston$medv01)
```



I will use the k-values 1,5,10

my three subset predictors will be all predictors,(nox, ptratio, lstat), and (tax,age,rm)

Question 4 (c):

#LOOCV

```
set.seed(1)
n <- nrow(newBoston)
boston.train <- cbind(newBoston$nox,
                      newBoston$ptratio,
                      newBoston$lstat)
boston.y.train <- newBoston$medv01

set.seed(1)
for(i in c(1,5,10)){
  boston.knn.cv <- knn.cv(train = boston.train, cl = boston.y.train, k=i)
  print(mean(boston.knn.cv != boston.y.train))
}

## [1] 0.229249
## [1] 0.173913
## [1] 0.1798419

set.seed(1)
n <- nrow(newBoston)
boston.train <- newBoston[, -14]
boston.y.train <- newBoston[, "medv01"]

set.seed(1)
for(i in c(1,5,10)){
  boston.knn.cv <- knn.cv(train = boston.train, cl = boston.y.train, k=i)
  print(mean(boston.knn.cv != boston.y.train))
}

## [1] 0.1936759
## [1] 0.1758893
## [1] 0.215415

set.seed(1)
n <- nrow(newBoston)
boston.train <- cbind(newBoston$tax,
                      newBoston$age,
                      newBoston$rm)
boston.y.train <- newBoston$medv01

set.seed(1)
for(i in c(1,5,10)){
  boston.knn.cv <- knn.cv(train = boston.train, cl = boston.y.train, k=i)
  print(mean(boston.knn.cv != boston.y.train))
}
```

```
## [1] 0.2549407
## [1] 0.2509881
## [1] 0.243083
```

The model with the predictor nox, ptratio, and lstat and k-value = 5 had the best results of all the models

Question 4 (d):

The variable in the Boston data set are not in the same scale. You can divide the data into training and test set and scale them

Question 4 (e):

```
set.seed(1)
n <- nrow(newBoston)
boston.train <- scale(cbind(newBoston$nox,
                             newBoston$ptratio,
                             newBoston$lstat))
boston.y.train <- newBoston$medv01

set.seed(1)
for(i in c(1,5,10)){
  boston.knn.cv <- knn.cv(train = boston.train, cl = boston.y.train, k=i)
  print(mean(boston.knn.cv != boston.y.train))
}

## [1] 0.1660079
## [1] 0.1778656
## [1] 0.1620553

set.seed(1)
n <- nrow(newBoston)
boston.train <- scale(newBoston[, -14])
boston.y.train <- newBoston[, "medv01"]

set.seed(1)
for(i in c(1,5,10)){
  boston.knn.cv <- knn.cv(train = boston.train, cl = boston.y.train, k=i)
  print(mean(boston.knn.cv != boston.y.train))
}

## [1] 0.1383399
## [1] 0.1264822
## [1] 0.1324111

set.seed(1)
n <- nrow(newBoston)
boston.train <- scale(cbind(newBoston$tax,
                             newBoston$age,
                             newBoston$rm))
boston.y.train <- newBoston$medv01
```

```

set.seed(1)
for(i in c(1,5,10)){
  boston.knn.cv <- knn.cv(train = boston.train,cl = boston.y.train,k=i)
  print(mean(boston.knn.cv != boston.y.train))
}

## [1] 0.2134387
## [1] 0.1758893
## [1] 0.1660079

```

The test error rate were lower for all the predictor subsets and k-values

Question 5 (a):

```

library(e1071)

## Warning: package 'e1071' was built under R version 4.1.3

mpg.median <- median(Auto$mpg)
newAuto <- data.frame(Auto)
newAuto$mpg01 <- ifelse(newAuto$mpg > mpg.median,1,0)
newAuto$mpg01 = as.factor(newAuto$mpg01)
newAuto$mpg = NULL

```

Question 5 (b):

```

set.seed(1)
auto.tune <- tune(method = svm,
                  mpg01~., data = newAuto,
                  kernel = "linear",
                  ranges = list(cost=c(0.001,0.01,0.1,1,5,10,100)))
summary(auto.tune)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.1
##
## - best performance: 0.08673077
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.13525641 0.05661708
## 2 1e-02 0.08923077 0.04698309
## 3 1e-01 0.08673077 0.04040897

```



```
## 4 1e+00 0.09961538 0.04923181
## 5 5e+00 0.11230769 0.05826857
## 6 1e+01 0.11237179 0.05701890
## 7 1e+02 0.11750000 0.06208951
```

Question 5 (c):

```
auto.svm <- svm(mpg01~., data = newAuto, kernel = "linear", cost = 0.1)
ypred <- predict(auto.svm,newAuto)
table(predict = ypred,truth = newAuto$mpg01)


##      truth
## predict 0    1
##      0 172    7
##      1  24 189

mean(ypred != newAuto$mpg01)

## [1] 0.07908163
```

Question 5 (d):

```
plot(auto.svm,data = newAuto, displacement~cylinders)
```



Question 6 (a):

```
library(ISLR)
set.seed(1)
n<- nrow(OJ)
oj.sample <- sample(1:n,800)
oj.train <- OJ[oj.sample,]
oj.test <- OJ[-oj.sample,]
```

Question 6 (b):

```
oj.svm <- svm(Purchase~.,
              data = oj.train,
              kernel = "linear",
              cost = 0.01)
summary(oj.svm)

##
## Call:
## svm(formula = Purchase ~ ., data = oj.train, kernel = "linear", cost =
## 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
## SVM-Kernel:  linear
##      cost:  0.01
##
## Number of Support Vectors:  435
##
## ( 219 216 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
```

There is 219 support vectors for class CH and 216 support vectors for class MM

Question 6 (c):

```
oj.train.pred <- predict(oj.svm,oj.train)
oj.test.pred <- predict(oj.svm,oj.test)
train.error <- mean(oj.train.pred != oj.train$Purchase)
test.error <- mean(oj.test.pred != oj.test$Purchase)
train.error

## [1] 0.175

test.error
```

```
## [1] 0.1777778
```

Question 6 (d):

```
set.seed(1)
oj.tune <- tune(method = svm,
                Purchase~.,
                data = oj.train,
                kernel = "linear",
                ranges = list(cost=c(0.001,0.01,0.1,1,5,10,100)))
summary(oj.tune)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.1
##
## - best performance: 0.1725
##
## - Detailed performance results:
##   cost   error dispersion
## 1 1e-03 0.31250 0.04124790
## 2 1e-02 0.17625 0.02853482
## 3 1e-01 0.17250 0.03162278
## 4 1e+00 0.17500 0.02946278
## 5 5e+00 0.17250 0.03162278
## 6 1e+01 0.17375 0.03197764
## 7 1e+02 0.17500 0.03486083
```

Question 6 (e):

```
oj.svm <- svm(Purchase~.,
              data = oj.train,
              kernel = "linear",
              cost = 0.1)
oj.train.pred <- predict(oj.svm,oj.train)
oj.test.pred <- predict(oj.svm,oj.test)
train.error <- mean(oj.train.pred != oj.train$Purchase)
test.error <- mean(oj.test.pred != oj.test$Purchase)
train.error

## [1] 0.165

test.error
```

```
## [1] 0.162963
```